

StitchUp

Shane T Fleming
Department of Electrical and
Electronic Engineering

Imperial College London
Email: shane.fleming06@imperial.ac.uk

David B Thomas
Department of Electrical and
Electronic Engineering

Imperial College London
Email: d.thomas1@imperial.ac.uk

Abstract—The abstract goes here.

I. INTRODUCTION

Aerospace and automotive systems often require strict constraints on performance along with a tight power budget making FPGAs an ideal candidate, yet they are seldom used. In such domains two main issues prevent widespread adoption: firstly they are difficult to program, with designs taking many man-hours to build and verify; and secondly they are highly susceptible to soft errors with single event upsets (SEU) potentially causing unwanted reconfigurations to the operating circuit. To address development time issues active research into High Level Synthesis (HLS) has received both academic and commercial success [4]. These tools aim to transform a widely known programming language, such as C, into a digital circuit removing the need for expensive hardware designers. As the issue of detecting/tolerating soft errors is typically done through replication and comparison costing large amounts of additional logic and power, which may require the design to break constraints.

This paper presents StitchUp an extension to the LegUp HLS tool [2] capable of producing circuits guaranteed to terminate even in the presence of SEUs provided that the input terminates, while always consuming equal or less logic than full circuit duplication. In order to achieve this the tool determines the set of instructions required to make an exact duplicate of the Control Flow Graph (CFG) while ignoring instructions which only effect the data flow. We argue that it is often the case that instructions which may influence the control flow of an application are more critical than instructions which never effect control; for example an image renderer which can tolerant errors, within it's pixel values but where errors in it's control flow could cause it to hang indefinitely [12].

The tool has three distinct stages: an analysis stage statically examines the input source and extracts any instruction which potentially effects control; a CFG shadow generation stage consists of a modified LegUp backend which takes the output of the analysis stage along with a description of the original schedule to produce a CFG shadow circuit; finally a wrapper generation stage combines the original circuit, the CFG shadow circuit, and generates error detection logic; Using this approach we are able to show that for certain applications we can guarantee that the CFG is followed correctly, while consuming only 4% of the resources required of standard Dual Modular Redundancy approaches.

Until recently soft error concerns have primarily been contained within domains where devices are placed in harsh

radioactive conditions, such as satellites in low earth orbit, however with shrinking device geometries this issue is set to become a problem for the entire industry[7]. If terrestrial devices start to become increasingly effected by ground level radiation then managing such faults is incredibly important especially in safety critical applications, such as driverless cars.

Protecting against soft errors is often achieved using N-Modular Redundancy, where N identical versions of the module are executed, their outputs are compared and any differences in the result indicate that a fault has occurred.

A. Background

Dual Modular Redundancy (DMR) is the traditional approach for detecting faults, where any difference between duplicate versions of operations on the same program state signals the precense of a fault. Typically this is performed with two identical hardware units executing in lock step and periodically comparing their outputs.

However redundancy is inherently costly consuming either power, or time, or often both, and is often not feasible in constrained systems such as satellites. For this reason a significant amount of research has been performed on attempting to reduce the amount of redundancy required to increase reliability. Due to the large body of work in this area this review will focus only on high-level approaches to reliability, examining current software techniques along with HLS techniques from the VLSI domain.

B. Generating Reliable Circuits using High Level Synthesis

Early VLSI work such as [1] examined how HLS tools could be used to reduce the amount of hardware required when duplicating the CDFG. The authors present a *Force-directed scheduling and alloction* algorithm, which aims to reduce the complete duplication of the resources by exploiting ones that are idle in each scheduling step.

Almost a decade later reliable HLS is revisited in [13], however instead of attempting to increase reliability through redundancy they increase reliability by using different versions of the components. They create a library characterisation of components which includes a reliability metric along with power, area, and latency. During the allocation phase of the HLS tool the library is searched to select the most reliable component while still meeting the other constraints. To achieve this a complete allocation, binding, and scheduling HLS algorithm is presented using ASAP scheduling. It starts by finding

the most reliable solution, then it minimises the number of resources required, and finally finds a valid schedule that meets the designs latency requirement. If constraints cannot be met a victim node is selected and swapped out for a different less reliable version, and keeps repeating until a feasible solution is found.

Two years later Glaß et al expanded on [13] in [5] where instead of attempting to search the library of components for the most reliable combination, they present the result to the user as a design space exploration. Along with selecting different types of hardware units, this work also takes into account the replication of various components while still respecting other constraints. To perform the multi-objective design space exploration, an evolutionary algorithm is used however the runtime to search the space can be quite costly and is not always able to find feasible solutions.

Recently work In [6] extends [13] so that along with the library characterizations for reliability the length of time that a particular FU is active is taken into account. They argue that the length of time that the circuit is active will effect the reliability metric; a full allocation, binding, and scheduling is found using an ILP solver.

Other recent work includes [3] where the vulnerability of variables is analysed and the most vulnerable variables for a HLS program are protected by assigning them to special registers. The proposed variable vulnerability analysis takes into account their lifetimes, dependencies, and branch probabilities. Through analysing the CDFG, error and branch probabilities are used to propagate error likelihood forward to estimate the total vulnerability for each variable in the design. A Integer Linear Program solver is then used to minimise the subset of protected registers to cover the maximum number of variables. They build their results on top of LegUp, and use LLVM to perform the variable vulnerability analysis. They achieve impressive results, where protecting 20% of the registers results in more than 60% soft error mitigation in the best case. However full fault injection tests are not performed on the generated hardware, and only state-based probabilistic error testing was performed.

Summarising, previous work on reliable HLS has generally revolved around the idea of adding a reliability metric to component libraries and using that information to build more reliable circuits. However, these approaches make no guarantees about reliability and only aim to improve overall systems reliability without paying much attention to how the reliability metric is derived. More recent work by Chen et al [3] is different in that a static analysis of the program is used to selectively duplicate sensitive variables.

C. Unreliable Hardware

Research into the use of unreliable hardware is slightly different to that of fault tolerance, since we can knowingly schedule instructions to execute on hardware that we don't trust. One tool that aims to aide the use of unreliable hardware is EnerJ [12] where type qualifiers are used to annotate whether a data type is approximate or precise. This tool takes an *all-or-nothing* approach to approximate computing, where precise annotated types are computed exactly, but imprecise ones

have zero guarantees. Strict constraints on how data flow occurs insure that the precise and approximate code regions are segregated; with approximate-to-precise data flow illegal, and precise-to-approximate legal. The authors provide a Java extension for the annotated type system, and many examples of code that can tolerate approximate regions demonstrating energy savings between 10-50%.

RAISE [10] is a recent project that aims to schedule instructions taking into account temporal and spatial vulnerabilities. They call the most likely time an instruction will receive and error the vulnerable period and this is calculated for every instruction. The tool then adjusts the instruction schedule so that the most vulnerable instructions have the least vulnerable period, minimising the overall likelihood of an error in the system. An example of this is when a multiply operation causes a pipeline stall which in turn causes a *load* or a *store* operation to get stuck at the decode stage of the processor pipeline; the longer this operation is stuck in the decode stage, the more likely it's opcode or address could be corrupted.

A metric called the Instruction Criticality Factor (ICF) is determined for every instruction through analysing the total number of dependent instructions that would be effected if it experienced a fault. This metric is then used to rearrange the execution order of the most critical instructions in order to increase the overall reliability of the program.

Summarising, the unreliable hardware field has also started looking towards high-level approaches to managing imprecise compute units. Either through a type system where users propose precise and imprecise regions of code, or through analysis of code and low level architecture models to influence scheduling decisions.

D. Software Soft-Error Protection

The hardware expense required for reliability has provided motivation for research into software fault tolerance. Many papers argue that with shrinking feature sizes and increasing use in power saving techniques such as DVFS reliability will be a concern for all processor markets, such as desktops, which might not be able to afford double or triple redundant execution cores [11].

Some software fault tolerance research requires modification to the underlying architecture especially when making use of multicore and multithreaded processors. One example of this is SRTR [14], simultaneously and redundantly threaded processors with recovery, which achieved fault tolerance through using two simultaneously running identical threads. However instead of running both threads in lock step one of them is executed slightly before the other. The trailing thread is responsible for checking that the execution of the leading thread is valid and completed before the leading thread commits it's results. To achieve this a new architecture is proposed which uses queued access to the register file in for checking results and to prevent it from becoming a bottle neck.

In an attempt to reduce architecture dependence research into software only fault tolerance techniques became popular in the mid 2000's. Early notable work include EDDI (Error Detection by Duplicating Instructions) by Oh et al [9] which

is a compiler tool for super-scalar processors that attempts to duplicate instructions using different registers and variables. It tries to schedule operations to make use of instruction level parallelism while still preserving its fault tolerant properties.

In an attempt to further reduce the overhead introduced by software fault tolerance Oh et al, [8] used signatures to ensure the correctness of the control flow graph (CFG) of a program. A static signature is assigned to every block and compared against a runtime signature during execution, flagging an error if there is a mismatch. This ensures that a transition between two basic blocks after a branch operation is a valid one according to the CFG. This method protects against transient faults in the program counter, and is often used as a method to detect malicious activity trying to alter the execution path of a particular program.

SWIFT [11], is another popular software only fault tolerance tool that aims to combine the methods presented in [9] and [8] for a single threaded application. Through transforming the code it is able to reclaim unused instruction level resources for duplicating instructions and inserting comparison points to check the replications. By combining these various techniques SWIFT is able to achieve a 14% speedup to other single threaded techniques which can achieve full fault tolerance coverage.

Summarising, software fault tolerance techniques tend to either fall into two broad categories; techniques that require small architecture changes, such as modifying the pipeline to allow efficient checking between multiple redundant threads [14]; and techniques that are purely software requiring no architecture changes. In a similar fashion to reliable HLS various approaches have been modified to trade off reliability and fault coverage for overhead. Compiler techniques such as static analysis of the code play an important role, especially in the scheduling of redundant operations and in deciding which portions of the program should be protected.

II. CONCLUSION

The conclusion goes here.

REFERENCES

- [1] A. Antola, V. Piuri, and M. Sami. High-level synthesis of data paths with concurrent error detection. In *Defect and Fault Tolerance in VLSI Systems, 1998. Proceedings., 1998 IEEE International Symposium on*, pages 292–300. IEEE, 1998.
- [2] A. Canis, J. Choi, M. Aldham, V. Zhang, A. Kammoona, J. H. Anderson, S. Brown, and T. Czajkowski. Legup: high-level synthesis for fpga-based processor/accelerator systems. In *Proceedings of the 19th ACM/SIGDA international symposium on Field programmable gate arrays*, pages 33–36. ACM, 2011.
- [3] L. Chen and M. Tahoori. Reliability-aware register binding for control-flow intensive designs. In *Proceedings of the The 51st Annual Design Automation Conference on Design Automation Conference*, pages 1–6. ACM, 2014.
- [4] T. Feist. Vivado design suite. *White Paper*, 2012.
- [5] M. Glaß, M. Lukasiewicz, T. Streichert, C. Haubelt, and J. Teich. Interactive presentation: Reliability-aware system synthesis. In *Proceedings of the conference on Design, automation and test in Europe*, pages 409–414. EDA Consortium, 2007.
- [6] Y. Hara-Azumi and H. Tomiyama. Cost-efficient scheduling in high-level synthesis for soft-error vulnerability mitigation. In *Quality Electronic Design (ISQED), 2013 14th International Symposium on*, pages 502–507. IEEE, 2013.
- [7] E. Ibe, H. Taniguchi, Y. Yahagi, K.-i. Shimbo, and T. Toba. Impact of scaling on neutron-induced soft error in srams from a 250 nm to a 22 nm design rule. *Electron Devices, IEEE Transactions on*, 57(7):1527–1538, 2010.
- [8] N. Oh, P. P. Shirvani, and E. J. McCluskey. Control-flow checking by software signatures. *Reliability, IEEE Transactions on*, 51(1):111–122, 2002.
- [9] N. Oh, P. P. Shirvani, and E. J. McCluskey. Error detection by duplicated instructions in super-scalar processors. *Reliability, IEEE Transactions on*, 51(1):63–75, 2002.
- [10] S. Rehman, M. Shafique, F. Kriebel, and J. Henkel. Raise: Reliability-aware instruction scheduling for unreliable hardware. In *Design Automation Conference (ASP-DAC), 2012 17th Asia and South Pacific*, pages 671–676. IEEE, 2012.
- [11] G. A. Reis, J. Chang, N. Vachharajani, R. Rangan, and D. I. August. Swift: Software implemented fault tolerance. In *Proceedings of the international symposium on Code generation and optimization*, pages 243–254. IEEE Computer Society, 2005.
- [12] A. Sampson, W. Dietl, E. Fortuna, D. Gnanapragasam, L. Ceze, and D. Grossman. Enerj: Approximate data types for safe and general low-power computation. In *ACM SIGPLAN Notices*, volume 46, pages 164–174. ACM, 2011.
- [13] S. Tosun, N. Mansouri, E. Arvas, M. Kandemir, and Y. Xie. Reliability-centric high-level synthesis. In *Design, Automation and Test in Europe, 2005. Proceedings*, pages 1258–1263. IEEE, 2005.
- [14] T. Vijaykumar, I. Pomeranz, and K. Cheng. Transient-fault recovery using simultaneous multithreading. In *ACM SIGARCH Computer Architecture News*, volume 30, pages 87–98. IEEE Computer Society, 2002.